

# ZFSLite: A Lightweight C++ Implementation of the ZFS File System

Tuan N. A. Vu<sup>1\*</sup> and Vi P. Nguyen<sup>1\*†</sup>

<sup>1\*</sup>School of Information and Communication Technology, Hanoi University of Science and Technology, 1 Dai Co Viet, Hanoi, Vietnam.

\*Corresponding author(s). E-mail(s): [tuan.vna2417212@sis.hust.edu.vn](mailto:tuan.vna2417212@sis.hust.edu.vn); [vi.np2417218@sis.hust.edu.vn](mailto:vi.np2417218@sis.hust.edu.vn);

<sup>†</sup>These authors contributed equally to this work.

## Abstract

Modern file systems increasingly rely on advanced features such as Copy-on-Write (CoW), cryptographic integrity verification, and atomic snapshots to ensure data reliability and simplify storage management, which is the reason why ZFS and Btrfs file systems are widely utilized in data-critical infrastructures. In this report, we present ZFSLite, a lightweight, userspace file system (FUSE) implemented in C++ that demonstrates the core principles of modern advanced file systems. ZFSLite operates over a virtual block device and provides a functional implementation of CoW semantics, Merkle tree-based data integrity verification using SHA-256, and a persistent reference-counted block allocator. Furthermore, it supports essential POSIX operations, including cross-directory renames, permissions, and atomic snapshots with efficient garbage collection. Evaluation of our prototype demonstrates functional correctness and practical performance, achieving approximately 5,000 IOPS and 20 MB/s throughput during CoW operations.

**Keywords:** ZFS, Copy-on-Write (CoW), FUSE, Merkle Tree, File System, C++

## 1 Introduction

The proliferation of data-intensive applications has driven the evolution of file systems from simple block allocators to complex storage managers that provide advanced

features such as data integrity verification, atomic snapshots, and transparent volume management. Modern file systems, such as ZFS and Btrfs, utilize Copy-on-Write (CoW) semantics to ensure transactional consistency and facilitate instantaneous snapshotting without data duplication. Furthermore, these systems employ Merkle trees with cryptographic hashes to detect and prevent silent data corruption, a critical requirement for enterprise storage environments.

In order to deeply study and understand these mechanisms, we implemented ZFS-Lite, a lightweight, userspace file system that isolates the core concepts found in advanced file systems. Developed in C++ using the Filesystem in Userspace (FUSE) framework, ZFSLite allowed us to practically examine the fundamental operations of CoW block allocation, Merkle tree-based integrity checking via SHA-256, and snapshot management over a virtual block device. By building these mechanisms from the ground up in userspace, we were able to directly investigate the complexities of modern file system design without the overhead of kernel-level development. In this report, we detail our implementation of ZFSLite, explore its on-disk data structures, and present an evaluation of its performance and functional correctness.

## 2 ZFS Architecture and On-Disk Layout

To fully comprehend the advanced mechanisms of ZFS, it is necessary to first understand the structural foundation upon which it operates. Traditional file systems like ext4 or FAT32 rely on static partition maps and linear block allocations. ZFS, however, integrates volume management and file system logic, operating on a dynamic tree-based structure.

### 2.1 Pooled Storage Paradigm

Traditional file systems require pre-allocated partitions of fixed sizes, leading to fragmented and inefficient storage utilization. When a partition is full, administrators must perform complex and error-prone operations such as resizing volumes or migrating data. ZFS eliminates traditional partitions by introducing the concept of *storage pools* (zpools). Physical storage devices are aggregated into a unified pool from which multiple file systems (datasets) can draw space dynamically.

### 2.2 On-Disk Data Structures

In ZFSLite, we model the foundational structures of ZFS using strictly packed C++ structures. These structures represent the exact binary layout written to the physical storage medium. The hierarchy consists of the Uberblock (the root of trust), dnodes (representing files/directories), dirents (directory entries), and block pointers.

```
1 #pragma pack(push, 1)
2
3 // The fundamental block pointer: contains the address and the checksum
4 struct zfs1_blkptr {
5     uint64_t blk_no;
6     uint8_t hash[32]; // SHA-256 cryptographic hash
7 };
8
9 // The dnode represents a file or a directory
```

```

10 struct zfsl_dnode {
11     uint32_t type;           // ZFSL_DNODE_FILE or ZFSL_DNODE_DIR
12     uint64_t size;           // File size or directory byte length
13     uint64_t mtime;          // Timestamps
14     uint64_t ctime;
15     uint64_t atime;
16     uint32_t uid;             // User ID
17     uint32_t gid;             // Group ID
18     uint32_t mode;            // Posix permissions
19     zfsl_blkptr direct[10];   // Pointers to direct data blocks
20     zfsl_blkptr indirect;     // Pointer to an array of indirect blocks
21 };
22
23 // The Uberblock is the absolute root of the file system
24 struct zfsl_uberblock {
25     uint64_t magic;           // Format identifier
26     uint64_t txg;             // Transaction Group (version counter)
27     uint64_t root_blk;        // Pointer to the root directory's dnode
28     uint8_t root_hash[32];    // SHA-256 hash of the root dnode
29     uint32_t snapshot_count;
30     zfsl_snapshot snapshots[16]; // Snapshot metadata array
31 };
32
33 #pragma pack(pop)

```

**Listing 1** Core on-disk structures defining the ZFSLite layout.

### 2.2.1 Structural Analysis

The most critical innovation in this layout is the `zfsl_blkptr`. In legacy file systems, a block pointer is merely a 32-bit or 64-bit integer representing an LBA (Logical Block Address). In ZFSLite, every single pointer is a composite structure containing both the LBA and a 32-byte SHA-256 hash. This design physically binds data location to data validation, forming the backbone of the Merkle tree architecture.

The `zfsl_uberblock` acts as the apex of this tree. It tracks the current Transaction Group (`txg`), which acts as an atomic sequence number for file system states.

## 2.3 Block Allocation and Virtual Device (VDev)

To persist these structures, ZFSLite interacts with a Virtual Device (VDev) layer that abstracts raw file I/O into strict 4 KB block operations.

Traditional file systems track free blocks using a bitmap (1 bit per block). However, to support  $O(1)$  snapshots where multiple file system states can share the exact same physical block, a single bit is insufficient. ZFSLite implements a **Reference Counted Block Allocator**. Instead of a bitmap, it utilizes an array of 8-bit unsigned integers. A value of 0 indicates a free block, 1 indicates a block exclusively owned by the live file system, and values greater than 1 indicate a block shared by the live system and historical snapshots.

## 3 Deep Dive: Core Mechanisms and Implementation

In this section, we analyze the defining mechanisms of ZFS—Copy-on-Write, Merkle Tree integrity, and Snapshots. We provide a detailed theoretical analysis of why these mechanisms are necessary in modern storage, followed by an in-depth examination of how they are implemented within the ZFSLite engine.

## 3.1 Mounting and Path Resolution

Before any high-level feature can operate, the file system must be able to translate human-readable paths (e.g., `/home/user/doc.txt`) into physical block addresses.

In ZFSLite, path resolution traverses the directory tree block by block. However, because of the strict Copy-on-Write requirement, we cannot merely find the target block; we must remember exactly how we got there. The `resolve_path` algorithm records a `PathNode` trace for every step of the traversal.

```
1 uint64_t ZFS::resolve_path(const std::string& path, zfs_l_dnode* out, std::vector<  
2     PathNode*> trace) {  
3     // 1. Read the root dnode from the Uberblock  
4     if (!read_dnode(uberblock.root_blk, out)) return 0;  
5  
6     // 2. Push the root onto the trace stack  
7     if (trace) trace->push_back({0, uberblock.root_blk, *out, 0, -1});  
8  
9     // 3. Tokenize path and traverse down the tree  
10    std::vector<std::string> parts = split_path(path);  
11    uint64_t cur = uberblock.root_blk;  
12  
13    for (const auto& part : parts) {  
14        // Search through the directory's dirent blocks  
15        // ... (directory parsing logic) ...  
16  
17        if (found) {  
18            uint64_t par = cur;  
19            cur = dirent->dnode_blk;  
20            read_dnode(cur, out);  
21  
22            // Record the parent, the child block, the exact dirent index, and the  
23            // dnode state  
24            if (trace) trace->push_back({par, cur, *out, dirent_blk_no, dirent_idx  
25                });  
26        }  
27    }  
28    return cur;  
29 }
```

**Listing 2** Path resolution and trace recording mechanism.

This `std::vector<PathNode>` trace operates as the breadcrumb trail that the CoW engine will follow backward. Without this trace, it would be impossible to efficiently locate the parent pointers that must be updated when a leaf node is relocated.

## 3.2 Copy-on-Write (CoW) Transactional Engine

### 3.2.1 Theoretical Analysis

The fundamental problem in legacy file system design is the "update-in-place" paradigm. When an application modifies a file, the file system overwrites the existing data blocks on the disk. If a power failure occurs halfway through this overwrite, the block contains a mixture of old and new data, leading to catastrophic corruption.

Journaling (Write-Ahead Logging) was introduced to mitigate this. Ext4, for example, writes intent to a journal before modifying the main disk area. While safe, this incurs significant performance penalties because data is written twice.

ZFS radically redefines this by enforcing strict Copy-on-Write (CoW). In ZFS, an active data block is **never** overwritten. When data is modified, a completely new, free block is allocated, and the new data is written there. Because the location of the data has changed, the parent node (the file's dnode) must be updated to point to this new block. However, modifying the parent dnode in-place is also forbidden by CoW rules. Therefore, the parent dnode must be written to a new block. This update forces the parent's parent (a directory dnode) to be written to a new block, triggering a cascading wave of allocations that propagates upward until it reaches the Uberblock at the root of the tree.

The transaction is only finalized when a single, atomic disk sector write updates the Uberblock to point to the new root. If the system crashes at any point prior to this final Uberblock write, the file system remains perfectly intact, pointing to the old data.

### 3.2.2 ZFSLite Implementation

In ZFSLite, the `cow_propagate` function is the engine that drives this mechanism. It takes the `PathNode` trace generated by `resolve_path` and processes it in reverse (from leaf to root).

```

1 void ZFS::cow_propagate(std::vector<PathNode>& trace, zfs1_dnode& leaf) {
2     zfs1_dnode cur = leaf;
3     std::vector<uint64_t> deferred_frees;
4
5     // Traverse the trace from bottom (leaf) to top (root)
6     for (int i = (int)trace.size() - 1; i >= 0; --i) {
7         PathNode& node = trace[i];
8
9         // 1. Allocate a completely new block and write the current node state
10        uint64_t new_blk = alloc->alloc_block();
11        write_dnode(new_blk, &cur);
12
13        // 2. Record the old block to be freed later
14        deferred_frees.push_back(node.dnode_blk);
15
16        // 3. If we reached the root, update the Uberblock and stop
17        if (i == 0) {
18            uberblock.root_blk = new_blk;
19            compute_sha256(&cur, sizeof(zfs1_dnode), uberblock.root_hash);
20            break;
21        }
22
23        // 4. Otherwise, update the parent's pointer to reference the new block
24        PathNode& parent = trace[i - 1];
25        uint64_t old_db = node.dirent_blk;
26        uint64_t new_db = alloc->alloc_block(); // CoW the parent's dirent block
27        too
28
29        // ... (read old_db, update dirent pointer to new_blk, write to new_db) ...
30
31        // Update the parent dnode to point to the new dirent block
32        for (int s = 0; s < 10; ++s) {
33            if (parent.dnode.direct[s].blk_no == old_db) {
34                compute_sha256(d_buf, VDEV_BLOCK_SIZE, parent.dnode.direct[s].hash);
35                ;
36                parent.dnode.direct[s].blk_no = new_db;
37                break;
38            }
39        }
40        deferred_frees.push_back(old_db);
41    }
42}

```

```

39     cur = parent.dnode; // Move up the tree
40 }
41
42 // 5. Critical Phase: Free old blocks ONLY AFTER the new tree is complete
43 for (uint64_t blk : deferred_frees) {
44     cow_free_block(blk);
45 }
46 }

```

**Listing 3** The Copy-on-Write propagation engine.


This implementation meticulously enforces consistency. A subtle but critical component of this logic is the `deferred_frees` vector (lines 3, 14, 43). During the upward traversal, superseded blocks are added to this list but are *not* immediately released to the allocator. If a block were freed immediately at step 2, the `alloc->alloc_block()` call at step 4 might recycle that exact block for a higher-level directory node. If a crash occurred immediately after, the old block would be overwritten, destroying the previous consistent state of the file system. By deferring all frees until the entire propagation chain is resolved, ZFSLite ensures perfect atomicity.

### 3.3 Merkle Tree Data Integrity

#### 3.3.1 Theoretical Analysis

As storage capacities grow into the petabyte scale, the statistical probability of silent data corruption (bit-rot) approaches certainty. Magnetic interference, cosmic rays, degraded flash cells, or firmware bugs can alter data without throwing any hardware error. Traditional file systems inherently trust the hardware; when an application requests a file, the file system reads the blocks and blindly passes them up, propagating corruption into application logic or backups.

ZFS introduces the concept of End-to-End Data Integrity. The entire file system is structured as a Merkle Tree (Figure 1). Every block of data is cryptographically hashed. However, storing the hash inside the block itself is insecure (if the block is corrupted or maliciously altered, the hash can be altered to match). Instead, ZFS stores the hash in the *parent* block pointer. The root of this tree (the Uberblock) provides the ultimate anchor of trust. By verifying the hash dynamically upon every read, ZFS can definitively prove whether the data retrieved from the physical disk is perfectly identical to the data that was originally written.



**Fig. 1** Structure of a Merkle Tree, illustrating how cryptographic hashes of data blocks ( $m$ ) are propagated up to the root node.

### 3.3.2 ZFSLite Implementation

ZFSLite utilizes the SHA-256 algorithm to implement this tree. Every `zfs1_blkptr` contains a 32-byte hash array. The cryptographic validation is deeply embedded in the read path.

```
1 int ZFS::read(const char *path, char *buf, size_t size, off_t offset) {
2     // ... (path resolution and boundary checks) ...
3
4     for (uint64_t b = start_b; b <= end_b; ++b) {
5         // 1. Locate the trusted hash from the parent dnode's pointer
6         uint64_t target_blk = dn.direct[b].blk_no;
7         uint8_t* expected_hash = dn.direct[b].hash;
8
9         // 2. Perform raw disk I/O
10        char blk[VDEV_BLOCK_SIZE];
11        vdev->read_block(target_blk, blk);
12
13        // 3. Dynamically compute the SHA-256 hash of the retrieved data
14        uint8_t computed_hash[32];
15        compute_sha256(blk, VDEV_BLOCK_SIZE, computed_hash);
16
17        // 4. Validate integrity
18        if (memcmp(computed_hash, expected_hash, 32) != 0) {
19            std::cerr << "SILENT CORRUPTION DETECTED: Hash mismatch on block " <<
20                target_blk << "\n";
21            return -EIO; // Return standard I/O error to the OS
22        }
23
24        // 5. If validation passes, copy data to the user buffer
25        memcpy(buf + buffer_offset, blk + block_offset, copy_length);
26    }
27    return bytes_read;
28 }
```

**Listing 4** Active Merkle tree integrity validation during read operations.

As demonstrated in Listing 4, the hardware is never trusted. After raw disk I/O is performed, the CPU computes the SHA-256 digest of the block buffer. The `memcmp` operation enforces absolute equality. If a single bit in the 4 KB block is flipped by the virtual device, the resulting hash will be wildly different, triggering the `-EIO` return code. This active detection prevents the silent propagation of corrupted data to the user application.

## 3.4 Atomic Snapshots and Garbage Collection

### 3.4.1 Theoretical Analysis

A snapshot is a frozen, read-only replica of a file system at a specific point in time. Legacy snapshot mechanisms (such as LVM snapshots) often operate at the block layer and require duplicating entire chunks of data, which degrades I/O performance and consumes significant storage space.

ZFS creates snapshots instantaneously in  $O(1)$  time with zero initial space overhead by leveraging the CoW engine. Because CoW dictates that active blocks are never overwritten, the historical state of the file system is naturally preserved on disk. Creating a snapshot simply requires saving the current Uberblock pointer and ensuring that the historical blocks are protected from the garbage collector, even if the live file system deletes or modifies those files later.



### 3.4.2 ZFSLite Implementation

To implement this, ZFSLite relies on its 8-bit reference counting allocator. Creating a snapshot requires recording the root and recursively incrementing the reference count for every active block.

```
1 int ZFS::take_snapshot() {
2     std::lock_guard<std::mutex> lock(zfs_mutex);
3
4     // 1. Record the current state of the Uberblock into the snapshot array
5     uint32_t idx = uberblock.snapshot_count++;
6     uberblock.snapshots[idx].txg = uberblock.txg;
7     uberblock.snapshots[idx].root_blk = uberblock.root_blk;
8     memcpy(uberblock.snapshots[idx].root_hash, uberblock.root_hash, 32);
9
10    // 2. Traverse the entire tree to increment reference counts
11    inc_tree_refcounts(uberblock.root_blk);
12
13    // 3. Persist the allocator and Uberblock state
14    alloc->save();
15    save_uberblock();
16    return 0;
17 }
18
19 void ZFS::inc_tree_refcounts(uint64_t dnode_blk) {
20     zfs_dnode dn;
21     read_dnode(dnode_blk, &dn);
22
23     // Increment the dnode block itself
24     alloc->inc_ref(dnode_blk);
25
26     if (dn.type == ZFSL_DNODE_FILE) {
27         // Increment all direct data blocks
28         for (int i = 0; i < 10; ++i) {
29             if (dn.direct[i].blk_no != 0) alloc->inc_ref(dn.direct[i].blk_no);
30         }
31         // ... (handle indirect blocks) ...
32     } else if (dn.type == ZFSL_DNODE_DIR) {
33         // Traverse directory and recursively increment children
34         // ...
35     }
36 }
```

**Listing 5** Snapshot creation and reference counting mechanism.

Garbage collection becomes an elegant byproduct of this architecture. When the live file system deletes a file, it calls `alloc->dec_ref()`. If a block is shared with a snapshot, its reference count will decrease (e.g., from 2 to 1), but the block will not be freed. The block is only returned to the free pool when all snapshots referencing it are deleted and its reference count hits strictly zero.

### 3.5 FUSE Integration Layer

To expose these mechanisms to the operating system, ZFSLite utilizes the FUSE (Filesystem in Userspace) framework. A translation layer maps POSIX system calls (`getattr`, `mkdir`, `read`, `write`, `unlink`) to the corresponding ZFSLite engine methods. Because the engine is not inherently thread-safe, a single global `std::mutex` aggressively serializes all FUSE callbacks to ensure transactional integrity across multi-threaded I/O operations.

## 4 Evaluation

We evaluated ZFSLite on a Windows platform utilizing WinFSP over a 512 MB virtual device image. The evaluation suite bypasses the FUSE layer to directly measure the core engine’s functional correctness and raw performance overhead.

### 4.1 Functional Correctness

The file system passed all functional correctness tests, including cross-directory atomic renames, deeply nested directory traversals, and full snapshot lifecycle management (creation, rollback, and deletion).

Crucially, the evaluation suite validates the Merkle tree via active corruption injection. The test script writes a file, computes its physical location on the virtual device, and directly manipulates the binary image file to flip a single byte in the data block. Upon requesting a read of the file via the engine, ZFSLite successfully intercepts the operation, prints an integrity mismatch warning, and returns an `-EIO` error code, fully mitigating the bit-rot simulation.

### 4.2 Performance Benchmark

To quantify the overhead introduced by the advanced mechanisms, we measured the throughput of 100 consecutive 4 KB block writes to a single file. Unlike a legacy file system, each of these writes forces ZFSLite to execute a complete CoW transaction: 1. Allocate a new data block and write the payload. 2. Compute the SHA-256 hash of the payload. 3. Propagate the structural changes upward, allocating new directory nodes and recomputing their respective hashes. 4. Persist the reference-counted allocator state to the device. 5. Commit the new Uberblock to disk via synchronous I/O.

Despite this heavy transactional workload, the benchmark achieved an approximate throughput of 20 MB/s, equating to roughly 5,000 IOPS. While significantly slower than native kernel-level file systems, this performance is highly practical for a userspace prototype and accurately reflects the cryptographic overhead of Merkle tree maintenance.

## 5 Conclusion

ZFS has revolutionized storage by addressing the fundamental flaws of traditional file systems, offering unparalleled data safety, dynamic pooled storage, and zerospace snapshots. By implementing these mechanics in userspace, ZFSLite distills these enterprise storage concepts into an accessible prototype.

Through its comprehensive evaluation, the ZFSLite prototype successfully achieves cryptographically secure file operations, robust directory overflow handling, and active, real-time corruption detection. By building these mechanisms entirely in C++, we gained practical insights into the complex choreography of Copy-on-Write propagation and reference-counted garbage collection.

Despite its success as a conceptual model, ZFSLite has notable limitations inherent to its architecture. The reliance on the FUSE framework introduces an inherent performance bottleneck, further exacerbated by a single global mutex

(`std::lock_guard<std::mutex> lock(zfs_mutex)` in `src/fuse_layer.cpp`) that aggressively serializes all I/O operations. Furthermore, ZFSLite lacks RAID-Z parity support for physical disk redundancy and omits advanced POSIX features such as extended attributes and symlinks. Addressing these concurrency bottlenecks and implementing multi-disk volume management represent promising avenues for future development.

## 6 Teamwork and Contributions

The ZFSLite project was collaboratively developed by the team. Each member made significant contributions to the project:

- **Tuan N. A. Vu:** Implementing ZFSLite, report and slides authoring.
- **Vi P. Nguyen:** Implementing ZFSLite, report and slides authoring.

## 7 References

- Figure [1](#) Source: <https://medium.com/@aannkkiittaa/understanding-fri-polynomial-commitments-scheme-7391da74c9d9>